BIG DATA ANALYTICS Name: Lovedeep Singh SID: 16103104

Hadoop Installation on Mac OS X

Wed, Aug 21, 2019

Install Homebrew and Cask

```
$ ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
$ brew install caskroom/cask/brew-cask
```

Install Java

```
$ brew update
$ brew cask install java
```

Configure SSH

In order to keep the safety of Hadoop remote administration as well as user sharing among Hadoop nodes, Hadoop requires SSH protocol. First, go to `System Preferences -> Sharing`, change `Allow access for`: **All Users**. Then open Terminal, input `ssh localhost`, if terminal returns `Last login: Sun Jul 2 16:57:36 2017`, which means that you have configured SSH Keys successfully before.

If you suffer the problem of `ssh: connect to host localhost port 22: Connection refused`, it happens since the remote login is closed.

```
$ sudo systemsetup -getremotelogin
Remote Login: off
```

You need to open port 22 in Mac OS X:

To verify if SSH Localhost is working check for files ~/.ssh/id_rsa and the ~/.ssh/id_rsa.pub files. If they don't exist generate the keys using below command

```
$ ssh-keygen -t rsa
```

Enable Remote Login: "System Preferences" -> "Sharing". Check "Remote Login" Authorize SSH Keys: To allow your system to accept login, we have to make it aware of the keys that will be used

```
$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

Test login.

```
$ ssh localhost
Last login: Fri Mar 6 20:30:53 2015
$ exit
```

Install Hadoop

First, install Hadoop via Homebrew: `brew install hadoop`, it will install the hadoop under `/usr/local/Cellar/hadoop`. Then, you need to modify the configuration files.

Go to `usr/local/Cellar/hadoop/2.8.0/libexec/etc/hadoop`, then open `hadoop-env.sh`

export HADOOP_OPTS="$HADOOP_OPTS -Djava.net.preferIPv4Stack=true"

change to

export HADOOP_OPTS="$HADOOP_OPTS -Djava.net.preferIPv4Stack=true -Djava.security.krb5.realm= -Djava.security.krb5.kdc="
export JAVA_HOME="/Library/Java/JavaVirtualMachines/jdk1.7.0_79.jdk/Contents/Home"

Then configure HDFS address and port number, open `core-site.xml`, input following content in `<configuration></configuration>` tag

```
<!-- Put site-specific property overrides in this file. -->
 <configuration>
    <property>
        <name>hadoop.tmp.dir</name>
        <value>/usr/local/Cellar/hadoop/hdfs/tmp</value>
        <description>A base for other temporary directories.</description>
    </property>
    <property>
        <name>fs.default.name</name>
        <value>hdfs://localhost:8020</value>
    </property>
</configuration>
```

Configure `jobtracker` address and port number in map-reduce, first `sudo cp mapred-site.xml.template mapred-site.xml` to make a copy of `mapred-site.xml`, and open `mapred-site.xml`, add

```
<configuration>
    <property>
        <name>mapred.job.tracker</name>
        <value>localhost:8021</value>
    </property>
</configuration>
```

Set HDFS default backup, the default value is 3, we should change to 1, open `hdfs-site.xml`, add

```
<configuration>
    <property>
        <name>dfs.replication</name>
        <value>1</value>
    </property>
</configuration>
```

Before running background program, we should format the installed HDFS first, executing command `hdfs namenode -format`, when terminal returns a long inforamtion like:

```
17/07/02 16:11:05 INFO namenode.NameNode: STARTUP_MSG:
/************************************************************
......
17/07/02 16:11:07 INFO namenode.NameNode: SHUTDOWN_MSG:
/************************************************************
SHUTDOWN_MSG: Shutting down NameNode at haodemacbook-
pro.local/192.168.1.4
************************************************************/
```

It means that we finish HDFS configuration, and Hadoop is ready to launch. Besides, maybe you will get a warning

```
$ ... WARN util.NativeCodeLoader: Unable to load native-
hadoop library for your platform... using builtin-java
classes where applicable
```

It happens since you are running on 64-bit system but Hadoop native library is based on 32-bit. This is not a big issue. If it appears, you can fixed by refering this link: here.

Launch Hadoop

Go to `/usr/local/Cellar/hadoop/2.8.0/sbin`, execute:

```
$ ./start-dfs.sh # start HDFS service
```

```
$ ./stop-dfs.sh # stop HDFS service
```

Ternimal will return the following information:

```
Starting namenodes on [localhost]
localhost: starting namenode, logging to /usr/local/Cellar/
hadoop/2.8.0/libexec/logs/hadoop-zhanghao-namenode-
HaodeMacBook-Pro.local.out
localhost: starting datanode, logging to /usr/local/Cellar/
hadoop/2.8.0/libexec/logs/hadoop-zhanghao-datanode-
HaodeMacBook-Pro.local.out
Starting secondary namenodes [0.0.0.0]
```

It means the local service launched successfully, then open Resource Manager in browser through the link `http://localhost:9870`, you can see the following page

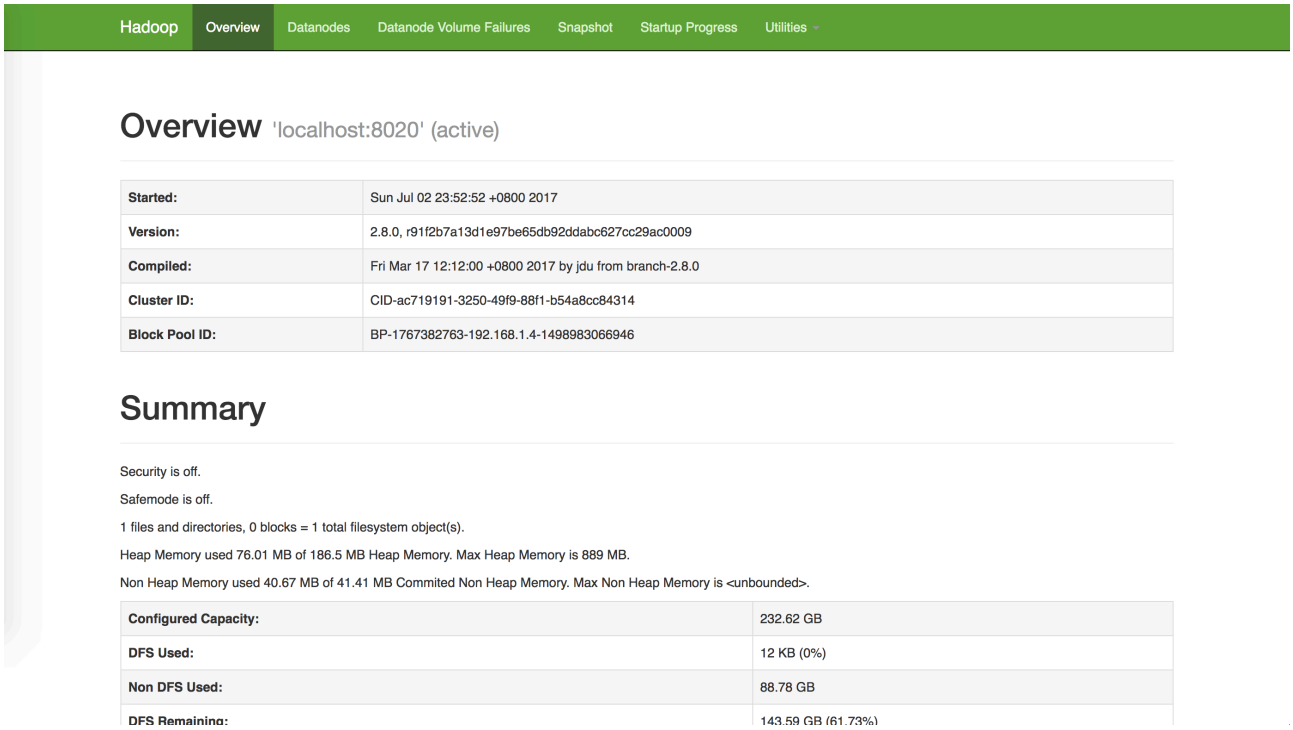

Note: earlier Hadoop versions of 2.x.x had 50070, after 3.0.0 it was changed to 9870 port number

Samely, under current diretory, you can start the JobTracker through the commands:

```
$ ./start-yarn.sh # start yarn, MapReduce framework
$ ./stop-yarn.sh # stop yarn
```

Then open browser and go to the page `http://localhost:8088`, Specific Node Information `http://localhost:8042`, you will see

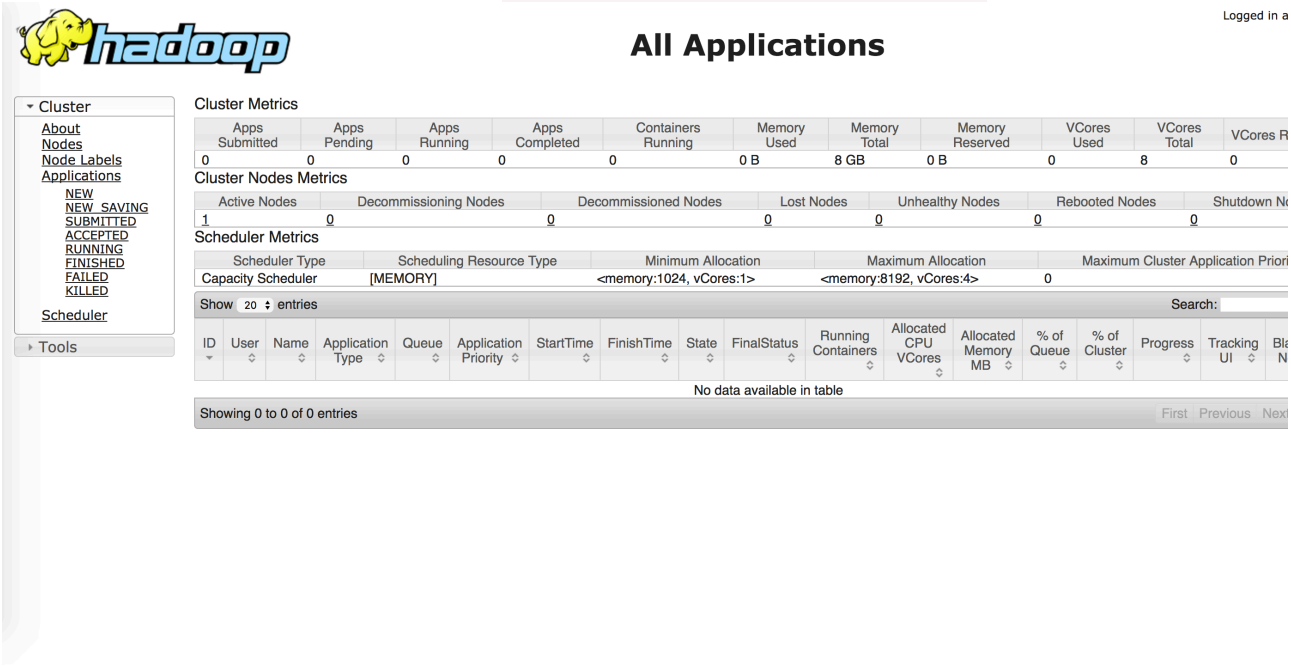

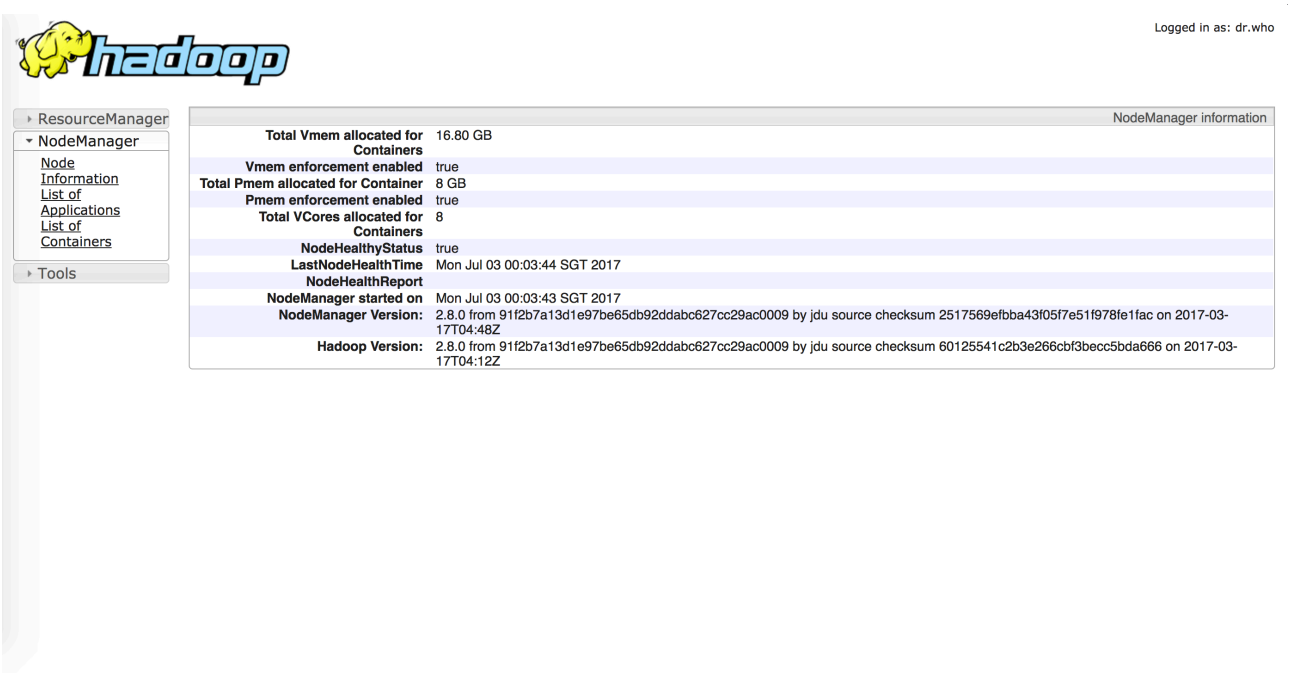

Simply, you can execute `./start-all.sh` and `./stop-all.sh` to start or close all the hadoop service. Finally, open `/etc/`

`profile` and add the configuration information of Hadoop environment variables.

```
export HADOOP_HOME=/usr/local/Cellar/hadoop/2.8.0
export PATH=$PATH:$HADOOP_HOME/sbin:$HADOOP_HOME/bin
```

Then you can start and close Hadoop under the user directory rather than go to `/usr/local/Cellar/hadoop/2.8.0/sbin` every time.

Note: Since 3.0.0 version, lot of port numbers changed. We have:

In fact, lots of others ports changed too. Look:

```
Namenode ports: 50470 --> 9871, 50070 --> 9870, 8020 --> 9820
Secondary NN ports: 50091 --> 9869, 50090 --> 9868
Datanode ports: 50020 --> 9867, 50010 --> 9866, 50475 --> 9865, 50075 --> 9864
```